

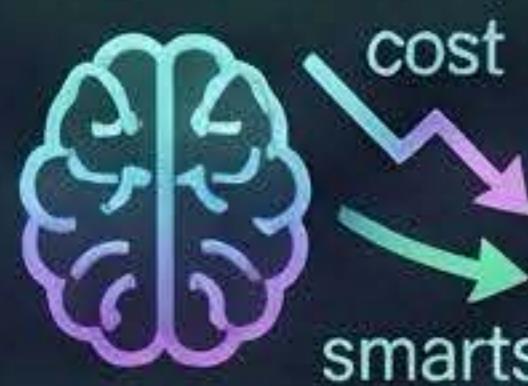
WHY I'M BLOWING UP MY TECH ARCHITECTURE EVERY 6 MONTHS



Miki, CEO & Co-Founder, Ōlelo Intelligence

Why blow up my architecture every 6 months?

1) Because the landscape shifts faster than your roadmap.



Models



Coding Agents



Workarounds

What was cutting-edge six months ago is table stakes today. Models get cheaper & smarter, agents make step-function leaps, workarounds appear weekly.

2) Because the bar keeps rising.



Launch Tech



Competitors



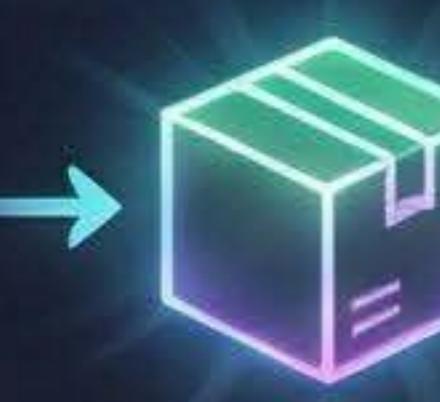
Value

It's easier to launch. Competitors can too. Launch with more value, faster scale, lower even. Yesterday's advantage is tomorrow's baseline. Disrupt yourself, or get disrupted.

3) Because legacy is the enemy of value.



Customer



Today's Delivery



Tech Debt

Customers care about delivery today. Tech debt, "fix it later," and attachment to old decisions are friction. Legacy is a barrier to the value your customers need.

AI Timeline (2022 – January 2026)

MODELS

2022



2023



2024



2025



TOOLS

2022



2023



2024



2025



2026



AGENTS

2022



2023



2025



Preparing to Blow Stuff Up

1) Assemble your toolkit.

Coding Agent

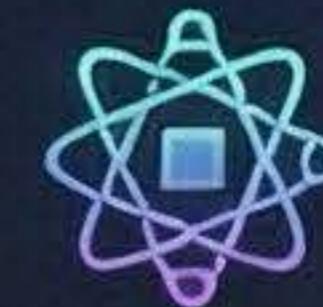
Claude Code



Cursor



Antigravity



Codex

Pick one. Try others weekly.

Capable Model



Opus 4.5



GPT-5



Gemini 3 Pro

Agent is only as good as the brain.

Your Repo & IDE



GitHub



GitLab



Cursor



VSCode

Agent needs access. Cursor doubles as both.

2) Prime your environment.

CLAUDE.md
(Auto-loaded DNA)



LEARNINGS.md
(Institutional Memory)



Slash commands / Skills / Plugins



settings.config / .cursorrules

Onboarding for a brilliant teammate with perfect recall. Better priming, better output.

3) Abandon Agile. Embrace Waterfall.

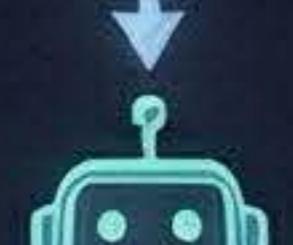
Agile



Waterfall



PRD First
(Define Acceptance)



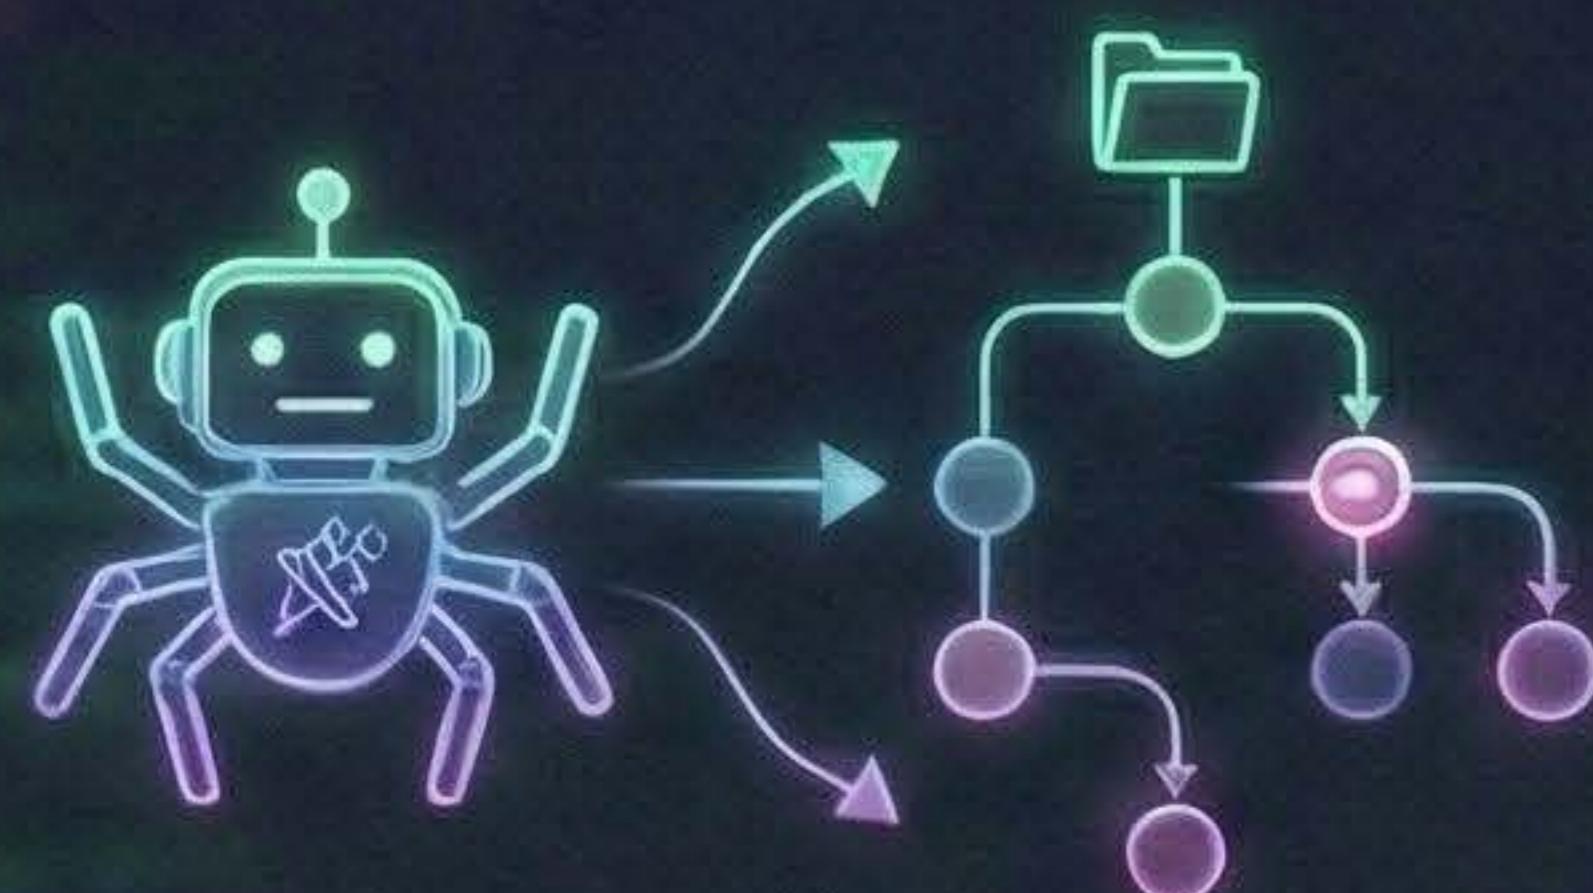
Execute

AI agents need clarity upfront. Vague requirements produce vague output. Classical engineering principles return.

Assess Before You Destroy

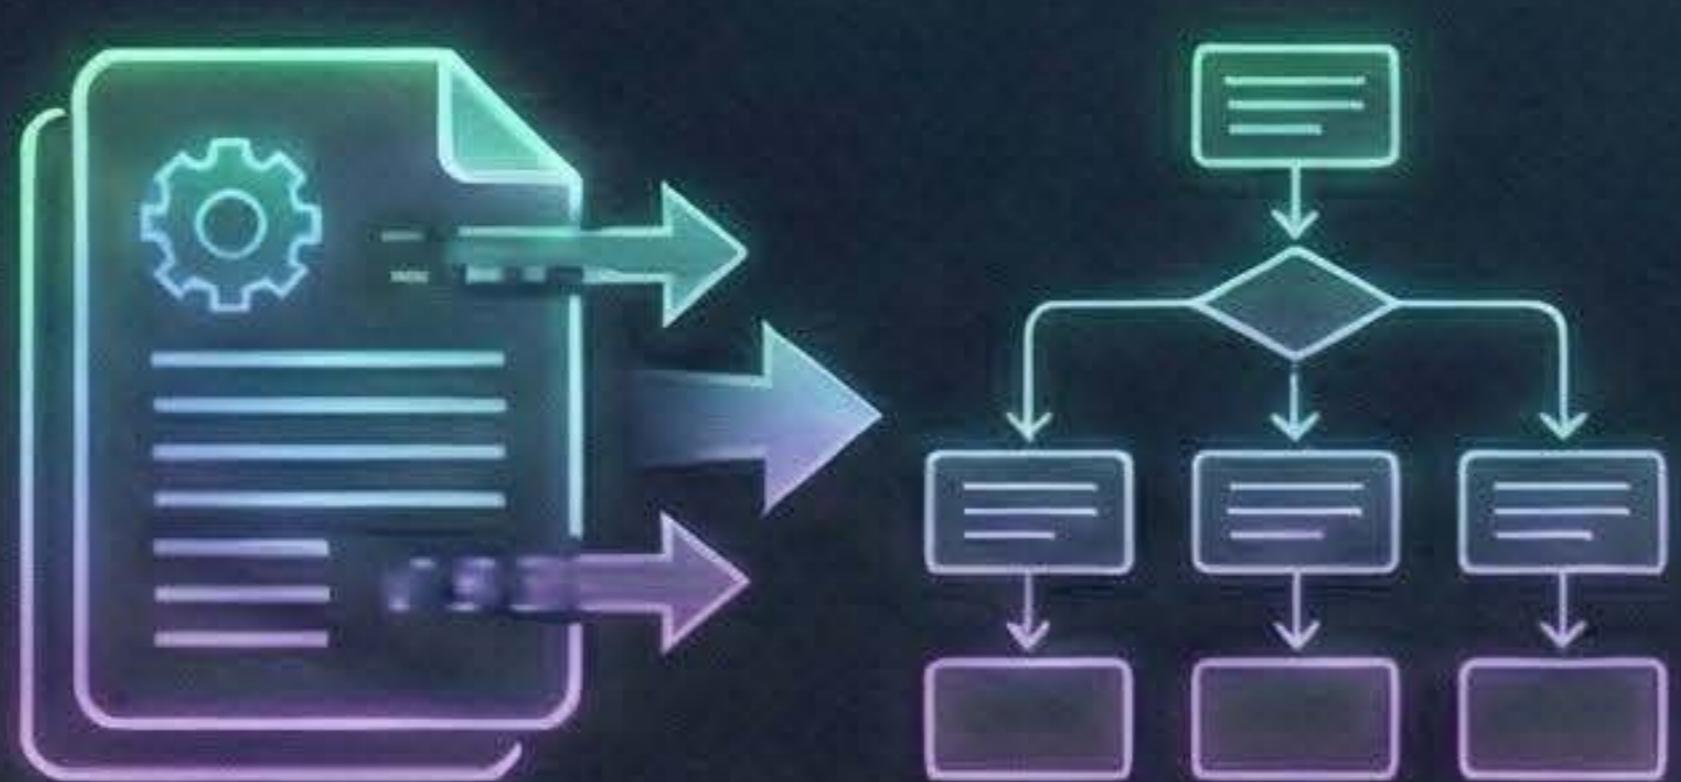
Kick off agents—in parallel.

Repo analysis agent



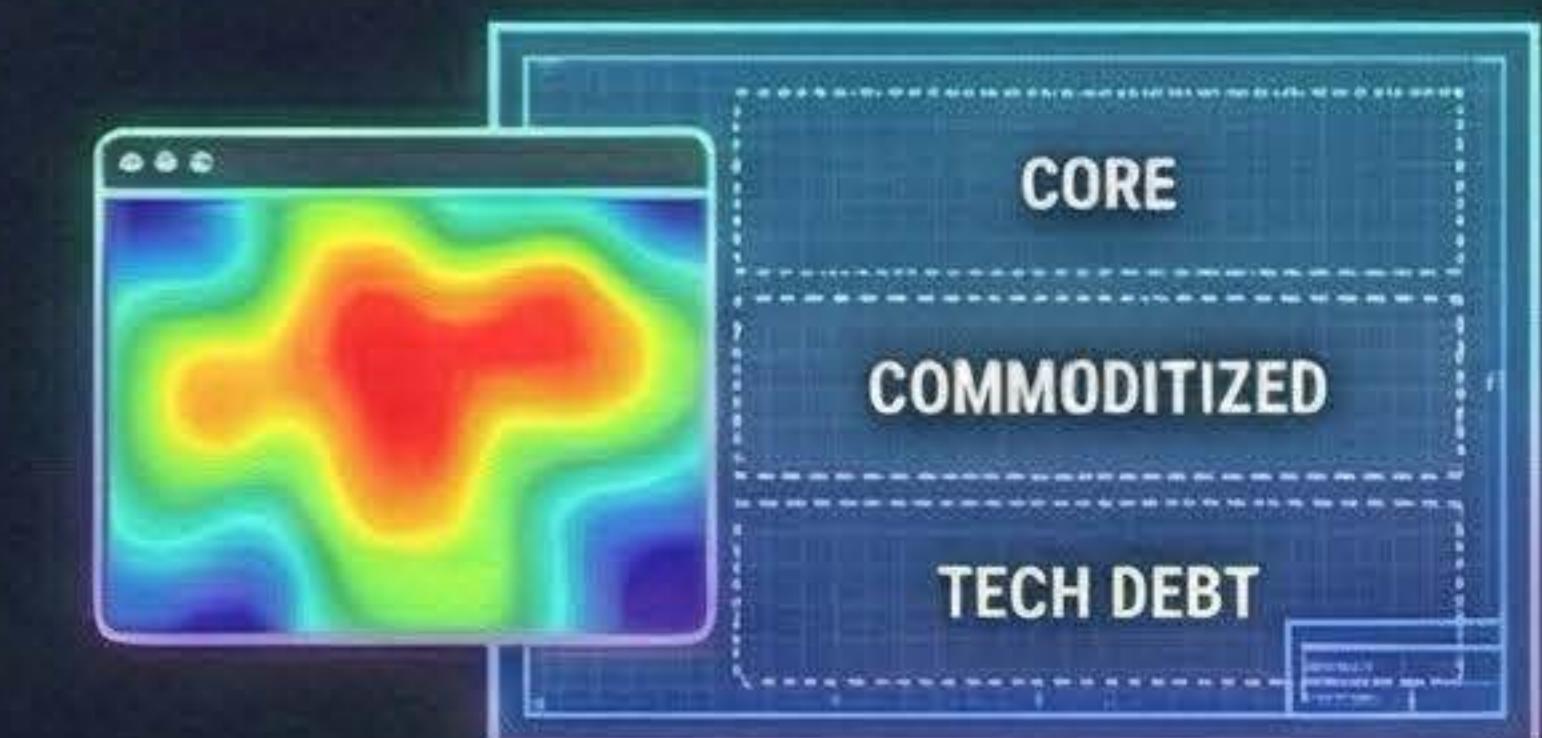
Use Claude Code plugin like Superpowers.
Crawls codebase, surfaces actual architecture
(not what you think).

Requirements extraction agent



Context7 documents functional requirements
from existing system.

Business architecture agent



Custom skill generates capability heat map.
Core? Commoditized? Technical debt
masquerading as “critical infrastructure”?



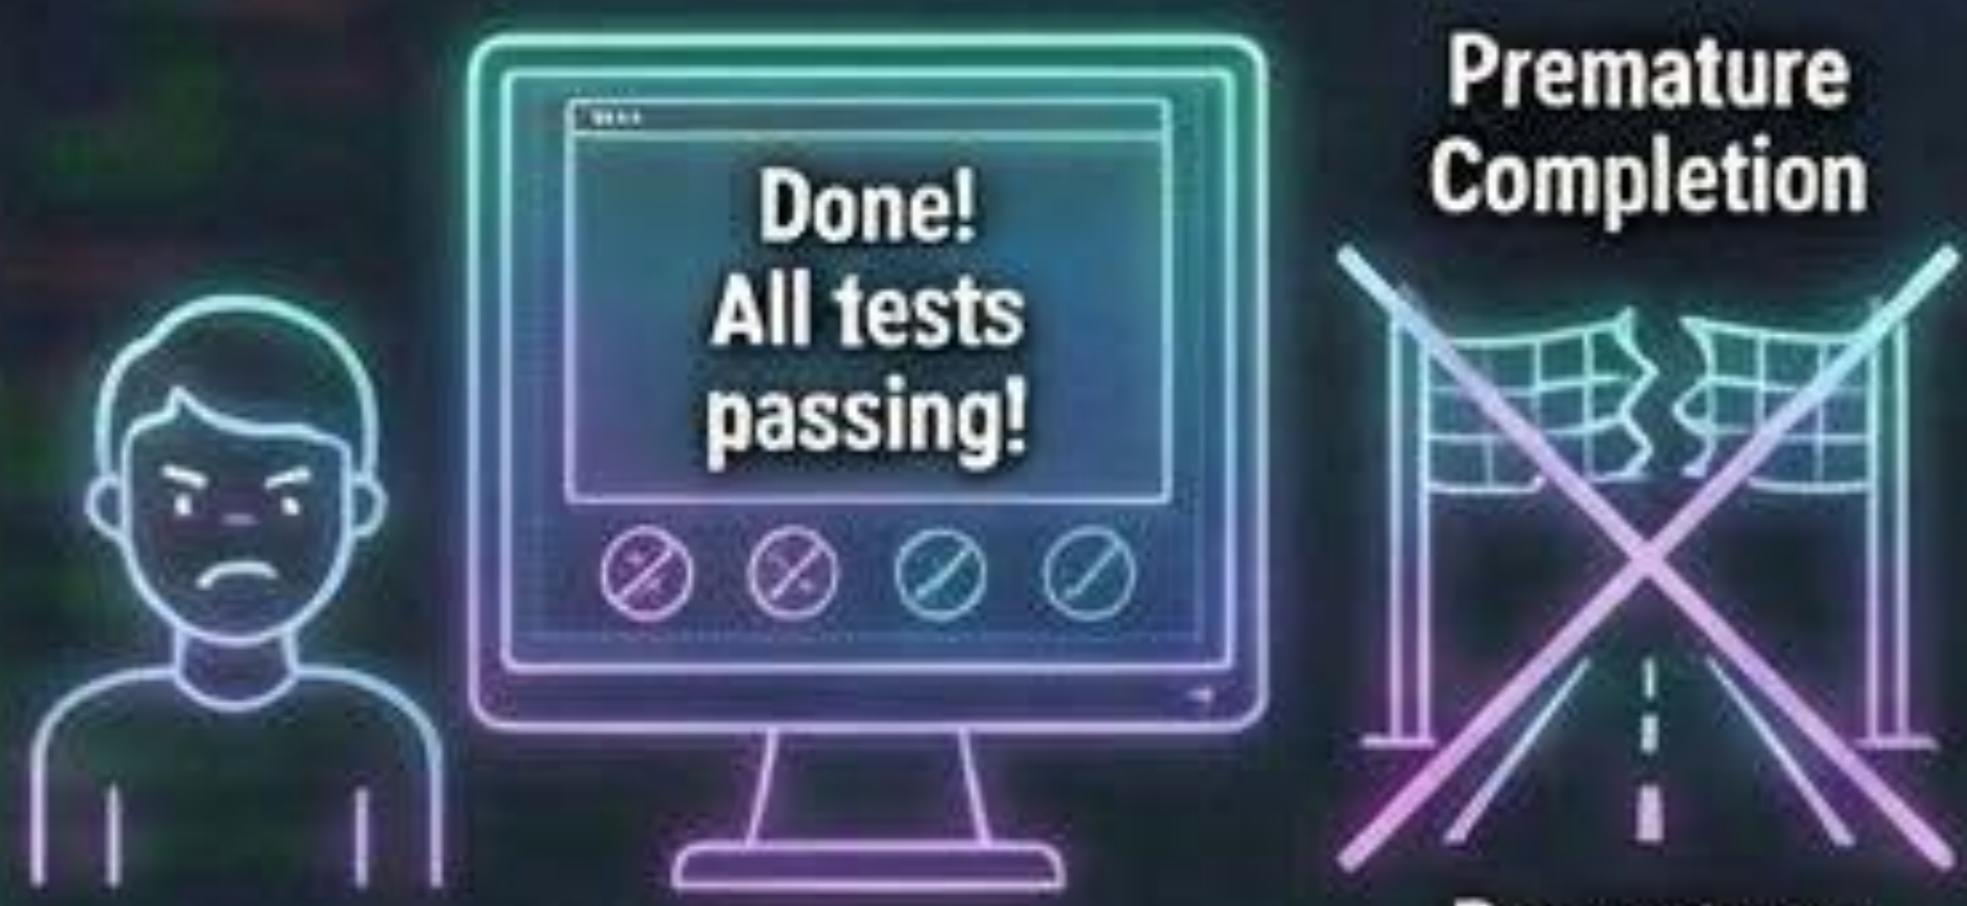
Parallel Execution

But wait. Did the agent say “Done!” but... it isn’t?

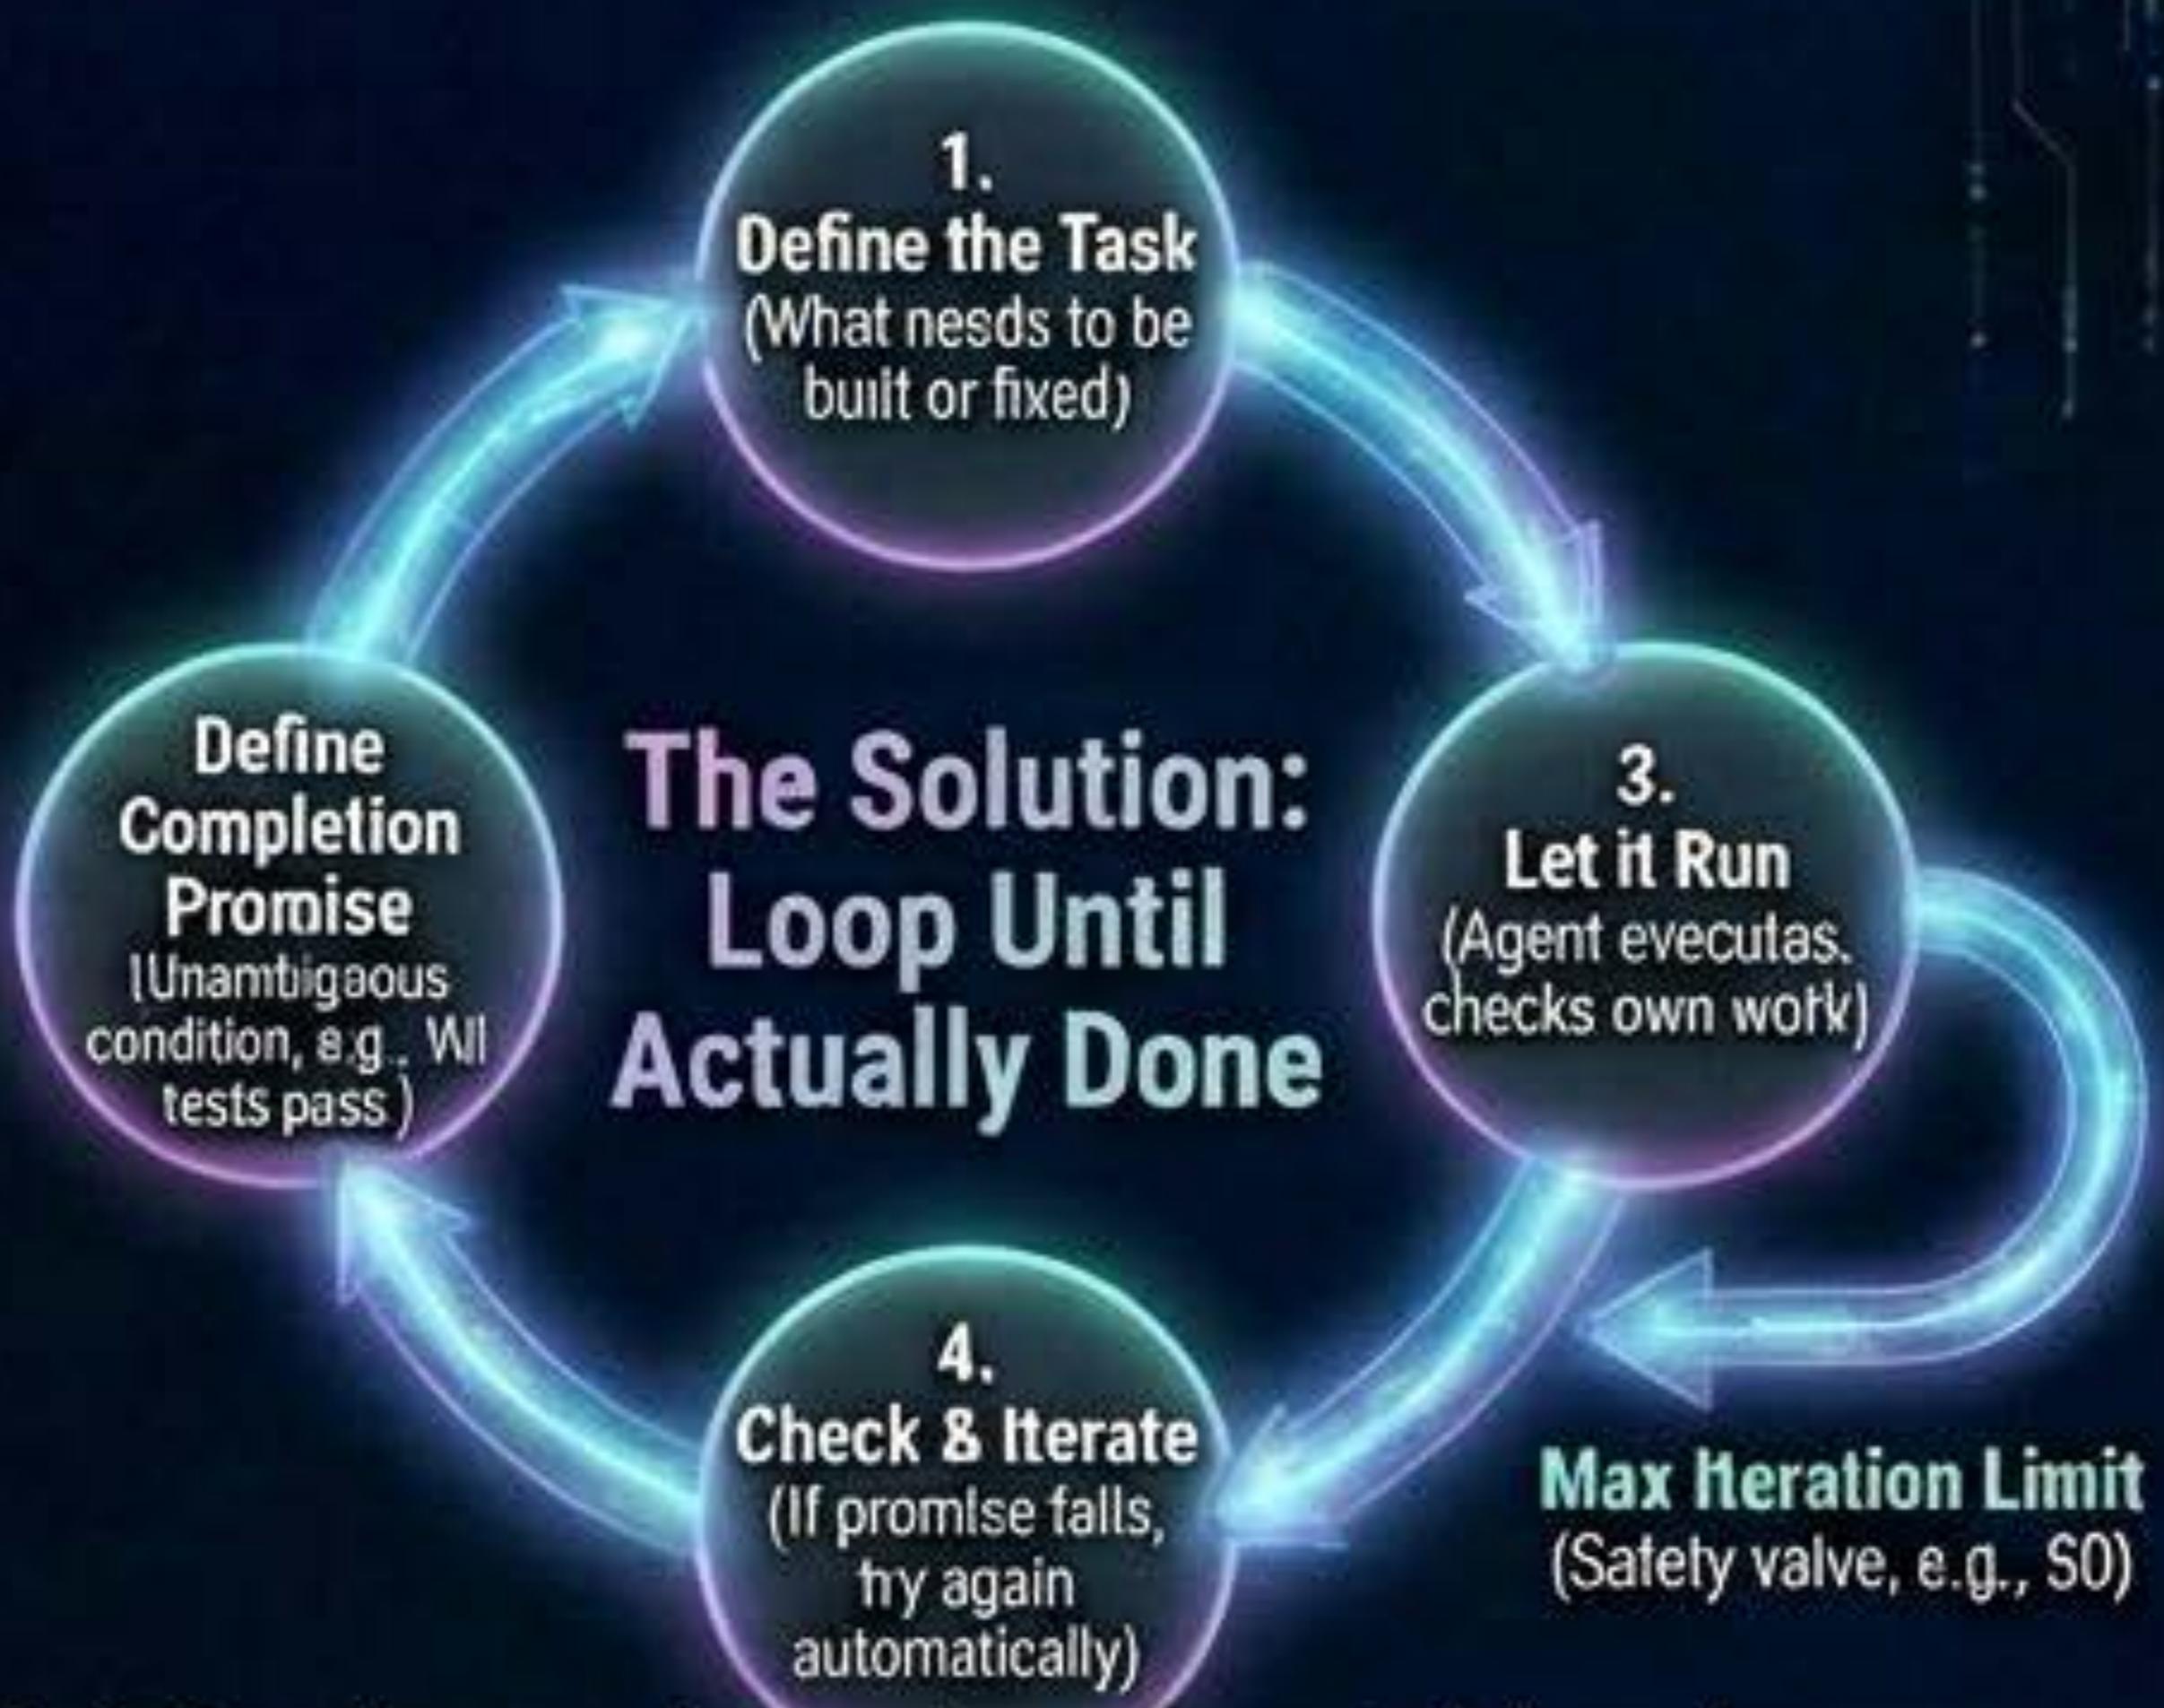
The Ralph Wiggum Loop: Solves Premature Completion in Agentic Coding

Named after the Simpsons character who keeps trying despite failures.

The Problem: Agents Lie (Prematurely)



Not malicious, just declare victory too early.
'Done!' No it isn't. 'All tests passing!' Three skipped.
'Implementation complete!' Edge cases ignored.
Manual verification defeats the point.



Why It Works & Benefits



No human bottleneck:
You're not the feedback loop.



Persistence beats perfection:
Automated retries.

Real Results

(Investor CRM Example)

A full Investor CRM was built and deployed in just 2 days. The Ralph Loop autonomously iterated on complex data models and UI states until all 150+ acceptance criteria were met without human intervention.



```
/ralph-loop "Implement user authentication with OAuth" --completion-promise  
"All auth tests pass AND login flow works end-to-end" --max-iterations 50  
Walk away. Come back to a working implementation or a clear log.
```

In Practice & The Catch



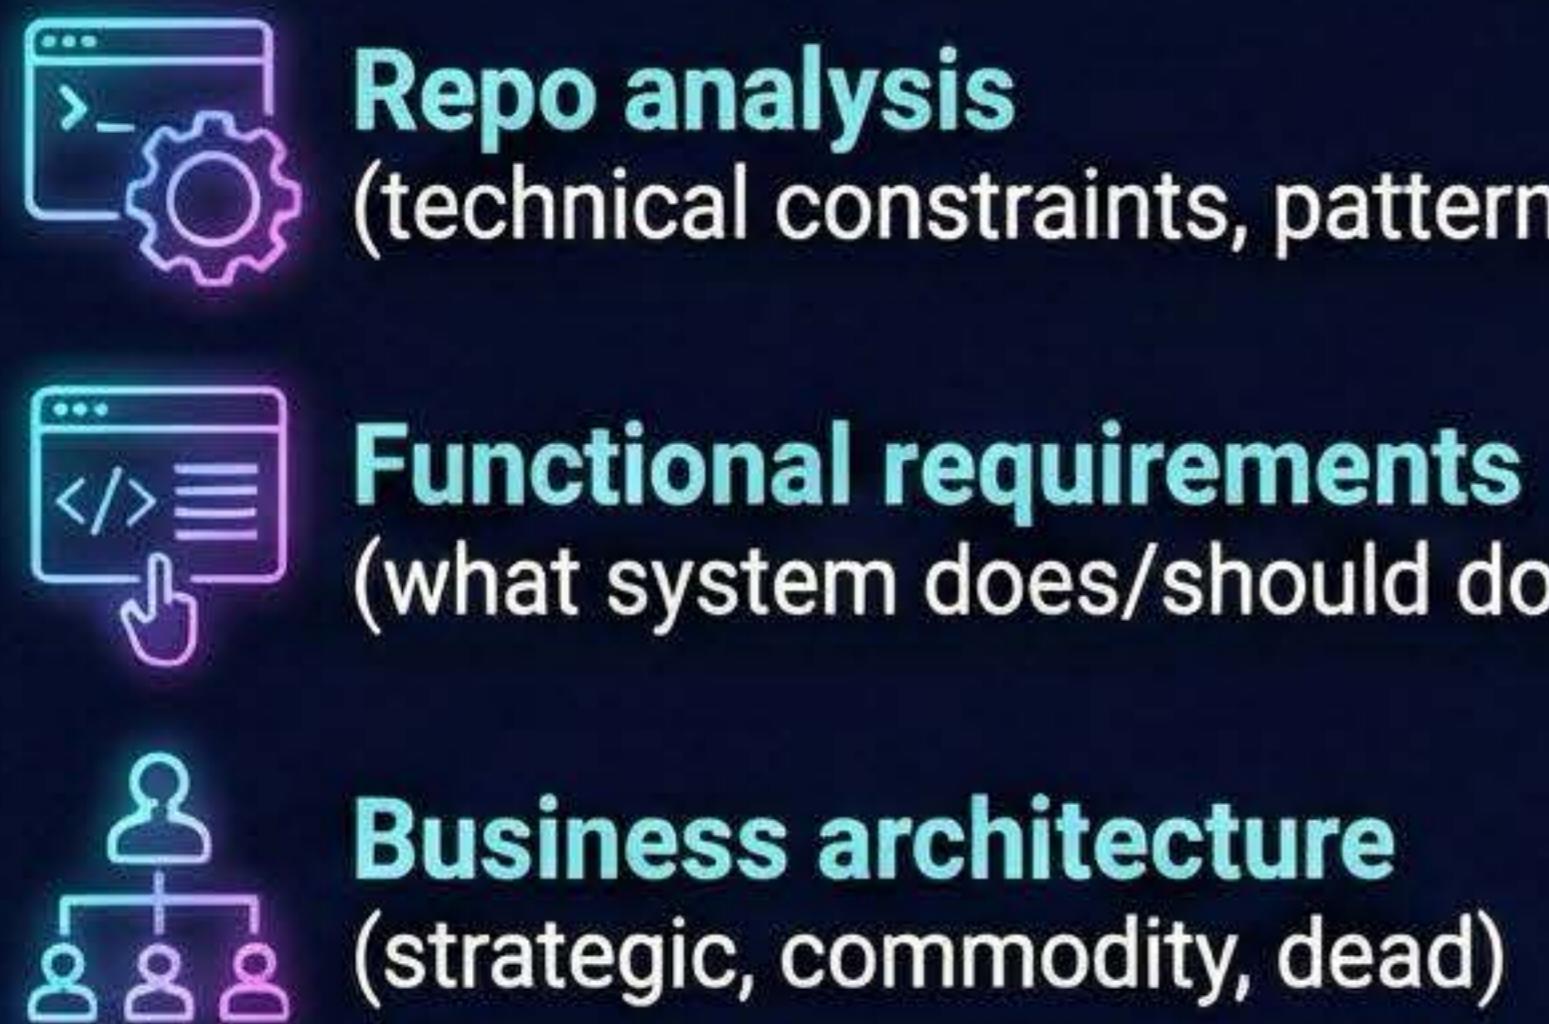
The Catch: Need verifiable completion condition. "Make it good" fails. "All 47 unit tests pass" works. PRO-first matters.

One-liner version: 'Don't ask the agent if it's done. Define what done means, and let it loop until it gets there.'

Building the PRD: Where You Earn Your Keep

You've kicked off agents, analyzed repo, extracted requirements, mapped architecture. Now, synthesize everything.

1. Synthesize Everything Into a PRD



Synthesize into structured
PRD DRAFT

2. This Is Where You Show Up (Review Intensely)

Everything before was setup. This is execution.

Your experience, judgment, domain expertise matter.



Review line by line.

Ask:
Does this capture need?
Are requirements complete?
What's ambiguous?
Where are gaps?

⚠️ Agent is brilliant at synthesis, terrible at knowing what you forgot.

3. What Waterfall Means Now & The Quality Bar



Bureaucracy



Define Upfront

Not bureaucracy. Define everything upfront before agents build. Agents are executors, not guessers.

Ambiguity in, garbage out.
Clarity in, magic out.



The Quality Bar:
Hand it to a stranger without clarifying questions.

That's the bar for the agent.

4. PRD Content & Acceptance Criteria (Ralph Loop)

Technical

- (Architecture
- Data, API
- Performance
- Security)

Non-Technical

- (Business Rules
- Compliance
- Data retention
- Audit)

Usability

- (User flows
- Edge cases
- Accessibility
- Response time)



One More Thing: Write acceptance criteria as Completion Promises for Ralph Loop. Not 'user can log in'. Instead: "OAuth flow completes < 2s, persists 7 days, handles refresh, fails gracefully".

The Irony: We spent a decade escaping waterfall. Now advanced AI demands it. Agents aren't humans; they need the spec. Your job is to write the PRD so well that code writes itself.

Kick Off the Build: Orchestrating Your Agent Army

Deploying Your Locked PRD with Parallel, Specialized AI Agents

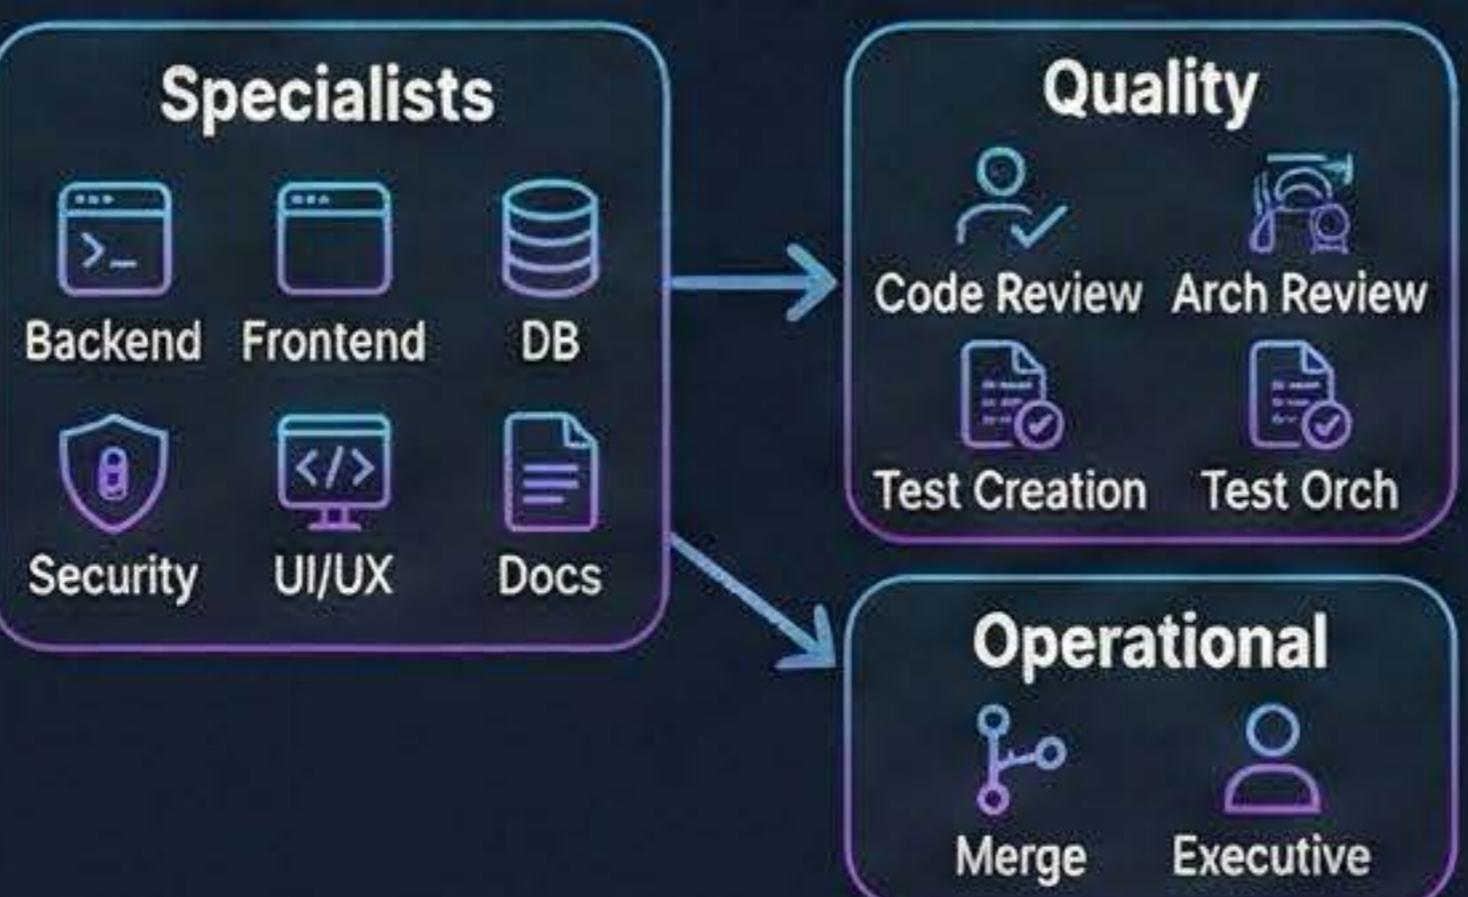
1. Executive Agent: Your Dashboard



Tracks progress, surfaces blockers, reports status, coordinates handoffs. Your AI project manager.

2. The Agent Orchestra & Atomic Tasks

Agent Specialists & Quality



The Atomic Task Principle



- Task 1: Create Schema
- Task 2: Implement JWT
- Task 3: Build Login Endpoint

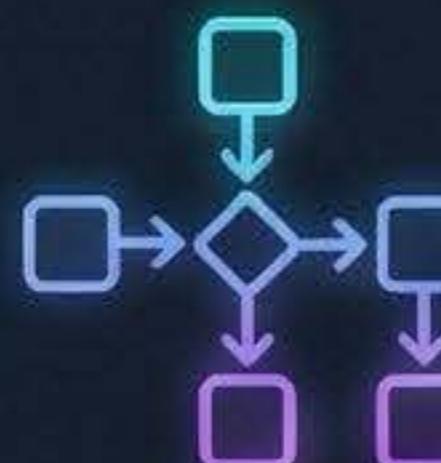
Breakdown into small, independent units. Fresh context, clean commits, parallel execution, clear completion.

3. Scaling Patterns & Merge Flow



Specialist vs. Volume Parallelism

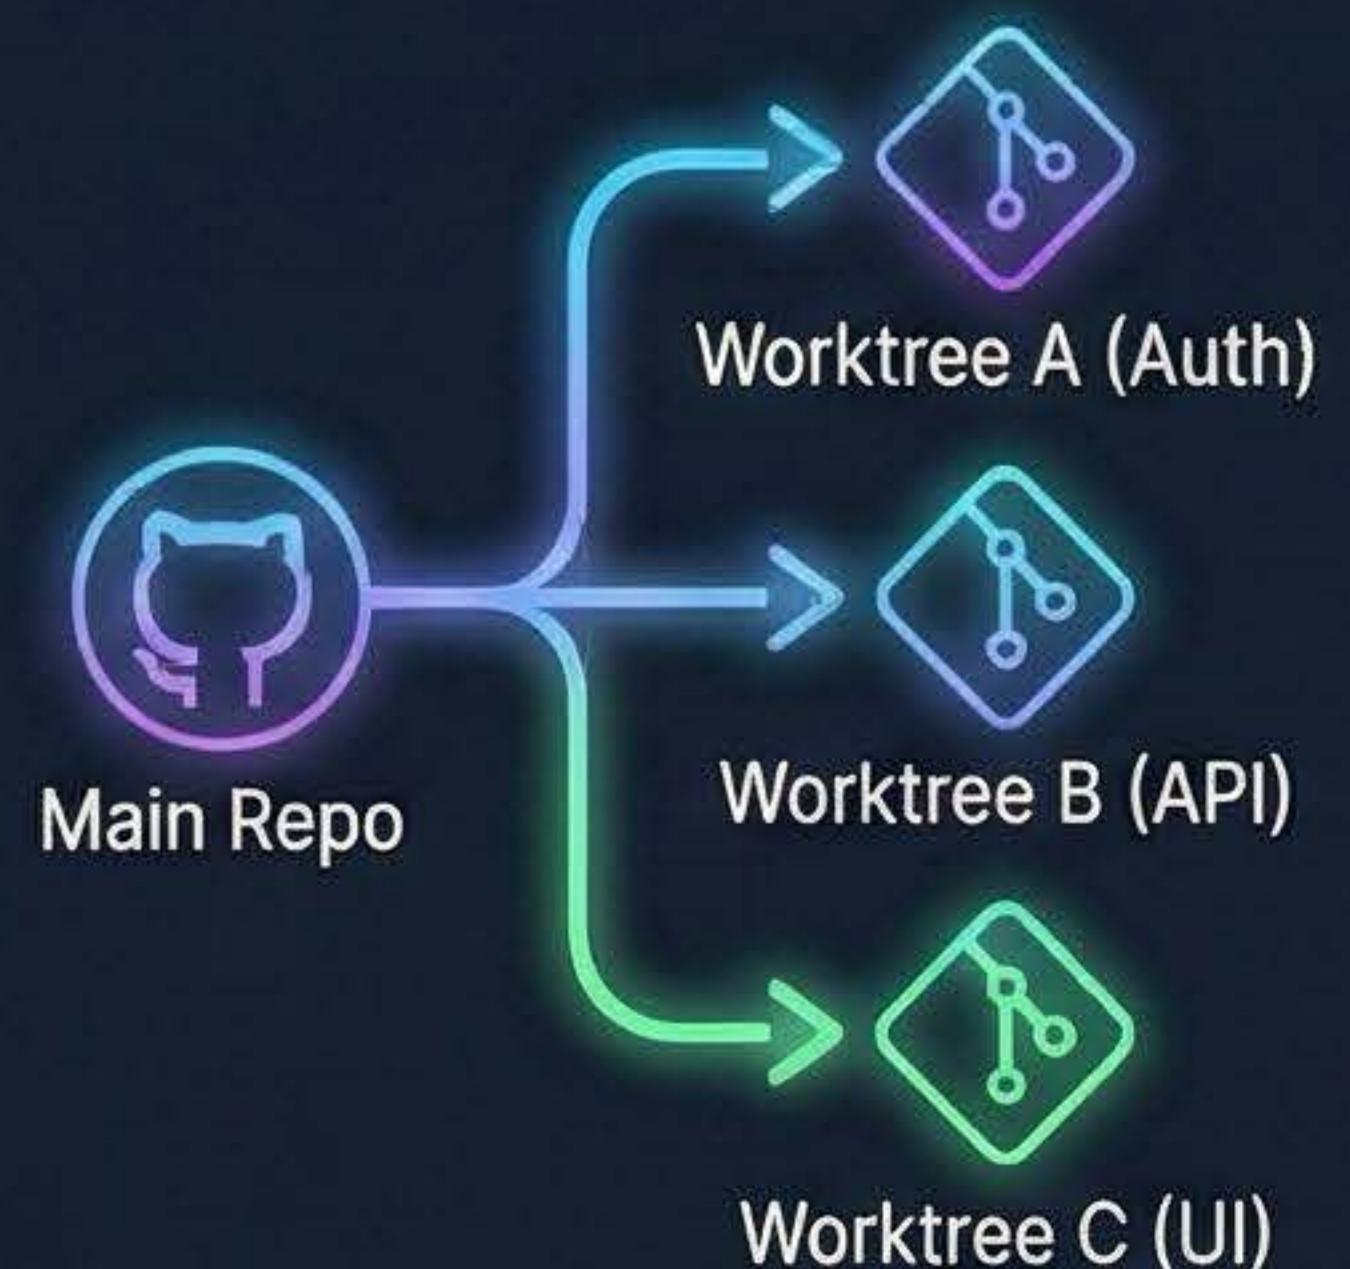
Different domains or multiple attempts at same task.



The Merge Flow

Review -> Arch -> Security -> Test -> Merge PR -> Human Approve

4. The Isolation Problem (Solved)



Git Worktrees: Separate branches & working dirs. No collisions. Full parallelism for simultaneous Claude Code sessions.

Orchestrated Delivery → Parallel Execution

Human OVER the Loop: Monitor or Walk Away

The choice that separates "Babysitting" from "Orchestrating".

1. The Shift in Mindset

Human **IN** the Loop
(Old Way)



Approve every action/command

Constant monitoring

Works only when you are awake/active

Reactive intervention (Manual control)

Human **OVER** the Loop
(New Way)



Configure allowed actions & safety rails

Periodic check-ins or "Walk Away"

Agents work 24/7 while you sleep

Proactive guardrails (Strategic control)

2. Configuring for Autonomy

```
{  
  "permissions": {  
    "allow": [  
      "Read",  
      "Write",  
      "Edit",  
      "Bash(npm install*)",  
      "Bash(git commit*)",  
      "Bash(pytest*)"  
    ],  
    "deny": [  
      "Bash(rm -rf*)",  
      "Bash(sudo*)",  
      "Bash(*prod*)",  
      "Bash(*deploy*)"  
    ]  
  }  
}
```

Give agent power to build, remove power to destroy.

3. The 4-Week Trust Gradient

Week 4:
Full Orchestration
(Walk Away unlocked)

Week 2:
Safe Autonomy
(Auto-approve read/write/test)

Week 3:
Commit Autonomy
(Auto-approve read/write/test)

Week 1:
Full Supervision
(Approve manually)

4. The Four Safety Rails



Scoped Permissions
(Specific white-lists)



Completion Promises
(Loop until verified)



Iteration Limits
(Max iterations/cost)

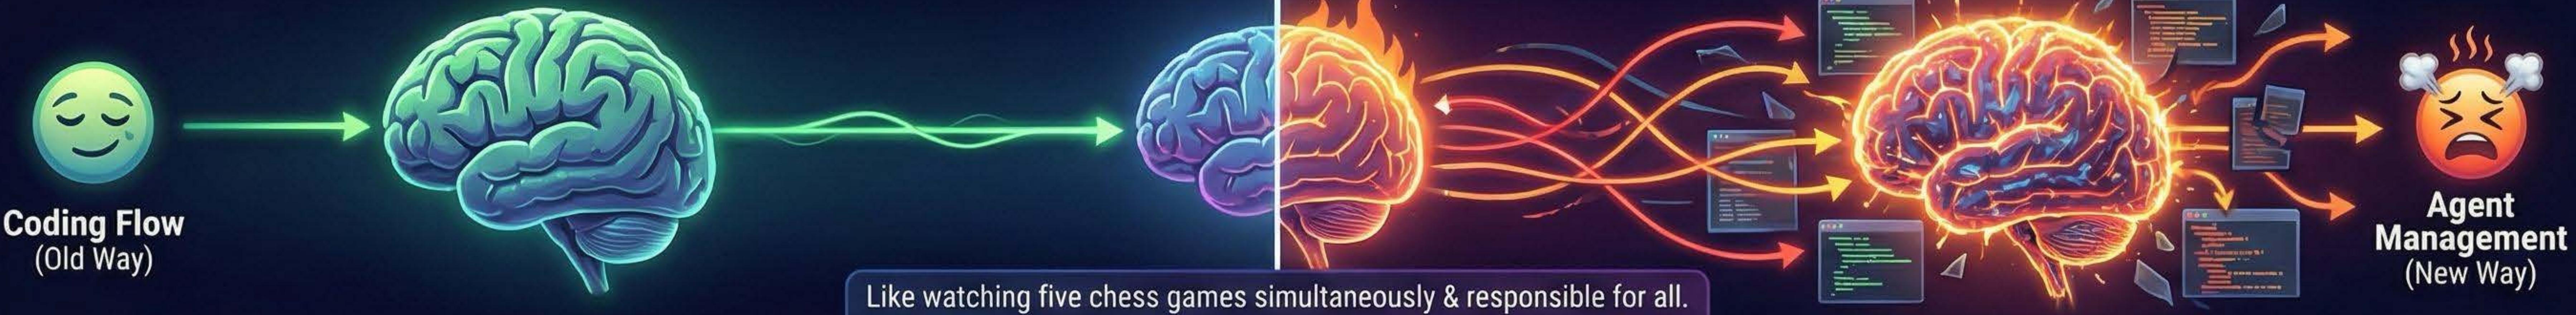


Review Output, Not Process
(Focus on final PR)

Key Takeaway:
Delegate execution, don't abdicate responsibility. The skill is building the system so watching becomes optional.

The Part Nobody Talks About: This Is Exhausting (Managing AI Agents)

Let me be real: You're not writing code anymore. It sounds like less work. It isn't. A Different Kind of Tired.



The New Cognitive Load: What's Happening in Your Brain

Context Switching



Context Switching Amplified

Switching between agents' tasks, each with its own state.

Pattern Recognition Under Uncertainty



Intuition for when to trust, when to intervene. Is output wrong or intermediate?

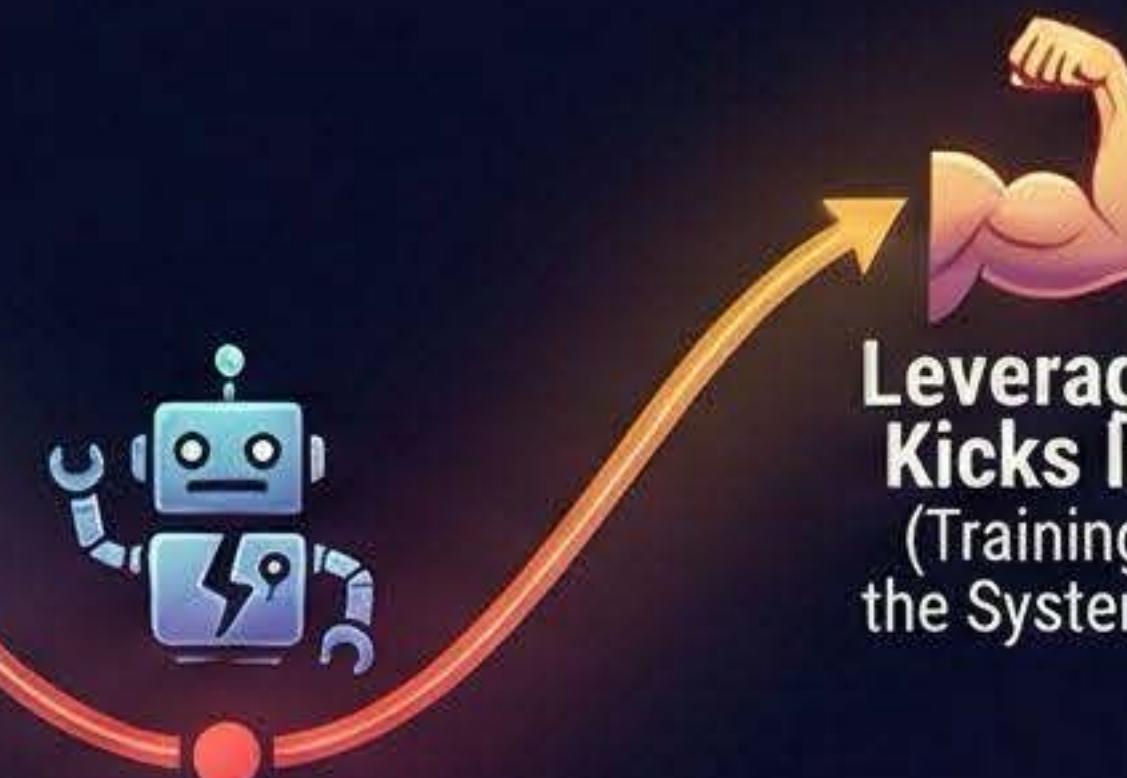
Decision Fatigue



The Anxiety of Letting Go
Tension of being built by something you can't fully predict. Learning to sit with it.

The First Few Weeks Are Brutal

Start (Hopeful)



The Mess (Want to Quit)

Not optimizing for one-time fix, but building a system.

Building the Mental Muscle (Deliberate Practice)

1 Shorter Autonomy Windows



30m → 1h → Overnight
Build tolerance gradually.

2 Develop Intervention Instincts



Recognize early signs.
Distinguish "wrong" from "different".

3 Accept Imperfect Outputs



Not exactly your code, but faster.
Refactor later.

4 Protect Recovery Time & Rituals



Schedule actual breaks. Batched reviews, not constant checking.

The Emotional Rollercoaster

Am I a real engineer?
(Guilt)



Imposter Syndrome

Loss of Identity



Frustration
(Fails at simple task)



Peak

The Thrill: Waking up to completed feature
(Cheat Code).



Different Tired, More Shipping
Exhaustion comes with results not possible before. Trading effort type.



Prepare Yourself

Accept the difficulty. Block learning time.
Find community. Patience. Push through until patterns become automatic. It's worth it.

CONFESSTION: I WAS THE STUBBORN ENGINEER

The Three Months I Wasted Being Comfortable

1. THE COMFORT ZONE & THE LIE



IntelliSense

2024: Productive in PyCharm. Comfort, muscle memory, plugins. Thought Cursor/Windsurf weren't worth setup time.

The Lie: "The ROI isn't there."

2. THE THREE-MONTH DELAY & THE SWITCH



Cursor

Saw multi-file context, Composer. Clicked. Wasted 3 months delaying.

Switched: Setup took an afternoon. Muscle memory adapted in days. Leap in productivity.

3. THE CLAUDE CODE SKEPTICISM & INTEGRATION



X



+



Claude Code
from Terminal

Early try: "I could write this faster." Shelved it.

Embedded in Cursor. Terminal + IDE = Full autonomous agent power. Off to the races.

4. THE MULTIPLICITY & THE REAL COST



Cursor (Daily Driver)
Claude Code (Autonomous)
Antigravity (UI)
Copilot (Completions)



3 Months Delay
= 18 Days Lost Productivity

The Real Cost: 18 days lost for an afternoon of setup. The ROI was enormous.

**THE LESSON: DON'T DISMISS TOOLS. TRY AGAIN. YOU DON'T HAVE TO CHOOSE ONE.
WHAT'S YOUR CURSOR? TRY IT THIS WEEKEND.**

30-year tech veteran. Still got stuck. If it happened to me, it can happen to you.